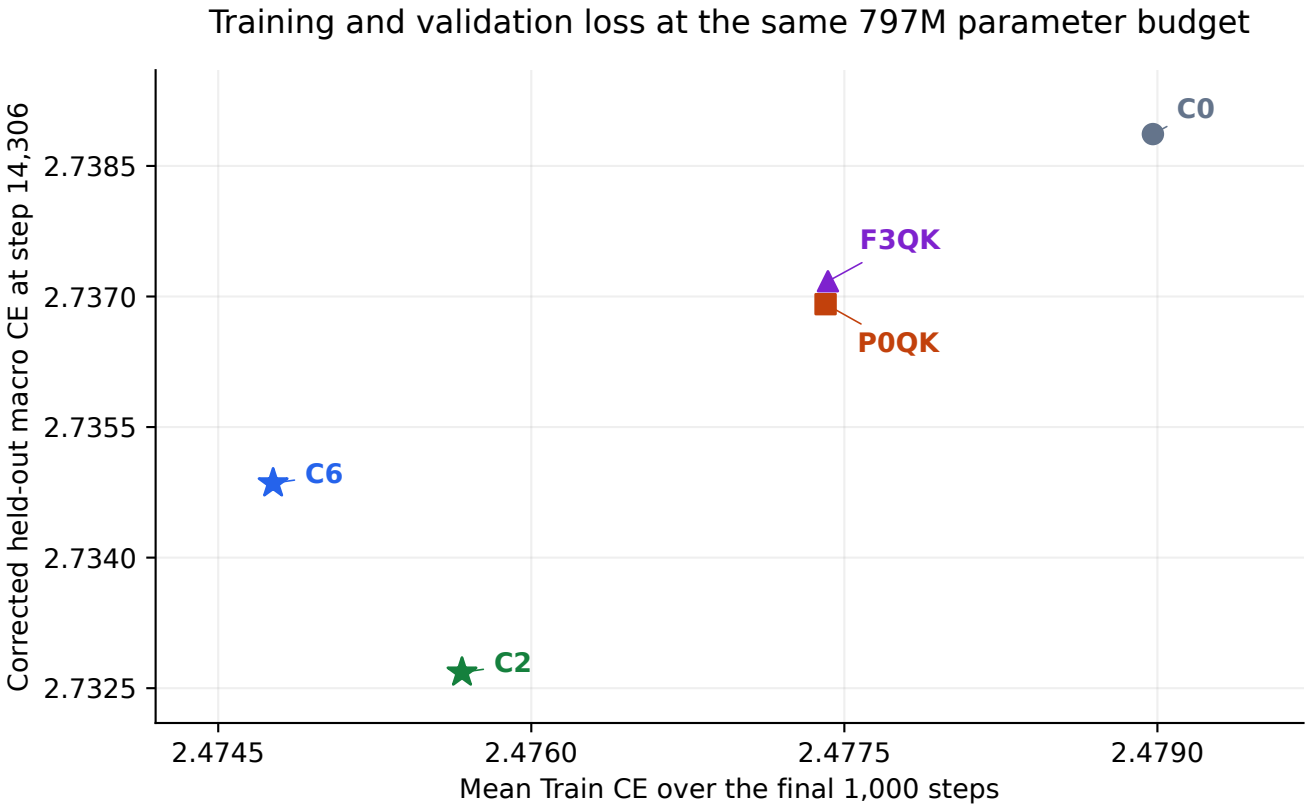
Training and validation loss at the same 797M parameter budget
Corrected held-out macro CE at step 14,306
2.7385
2.7370
2.7355
2.7340
2.7325
C0
F3QK
P0QK
C6
C2
2.4745
2.4760
2.4775
2.4790
Mean Train CE over the final 1,000 steps
Lower is better on both axes. Single seed; 30B tokens; 11 held-out sources.
Training allocations differ. Seed replication is pending; no significance claim.